

Amazon_Product_Ratings_&_Purchase_Behavior_Analysis

This project analyzes Amazon product data to explore how ratings, reviews, and other features relate to product purchases.

Import libraries

We import all required libraries for data processing, visualization, and analysis.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from statsmodels.stats.outliers_influence import variance_inflation_factor
from statsmodels.tools.tools import add_constant
import seaborn as sns
import statsmodels.api as sm
```

Load dataset

Load the dataset and preview the structure to understand available features.

```
data_url = "amazon_products.csv"
df=pd.read_csv(data_url)
```

```
df.head()
```

	product_title	product_rating	total_reviews	purchase_date_month	discount_percentage	original_price	is_best_seller	is_sponsored	has_coupon	buy_box_availability	delivery_date	sustainability_tags	product_image_url	product_page_url	data_collection_date	product_category	discount_percentage
0	BOYA BOYALINK 2 Wireless Lavalier Microphone f...	4.6	375.0	300.0	89.68	159.00	No Badge	Sponsored	Save 15% with coupon	Add to cart	2025-09-01	Carbon impact	https://m.media-amazon.com/images/I/71pAqiVEs3...	https://www.amazon.com/sspa/click?ie=UTF8&spc=...	2025-08-21 11:14:29	Phones	43.60
1	LISEN USB C to Lightning Cable, 240W 4 in 1 Ch...	4.3	2457.0	6000.0	9.99	15.99	No Badge	Sponsored	No Coupon	Add to cart	2025-08-29	NaN	https://m.media-amazon.com/images/I/61nbF6aVIP...	https://www.amazon.com/sspa/click?ie=UTF8&spc=...	2025-08-21 11:14:29	Laptops	37.52
2	DJI Mic 2 (2 TX + 1 RX + Charging Case), Wirel...	4.6	3044.0	2000.0	314.00	349.00	No Badge	Sponsored	No Coupon	Add to cart	2025-09-01	NaN	https://m.media-amazon.com/images/I/61h78MEXoj...	https://www.amazon.com/sspa/click?ie=UTF8&spc=...	2025-08-21 11:14:29	Laptops	10.03
3	Apple AirPods Pro 2 Wireless Earbuds, Active N...	4.6	35882.0	10000.0	162.24	162.24	Best Seller	Organic	No Coupon	NaN	NaN	NaN	https://m.media-amazon.com/images/I/61SUj2aKoE...	https://www.amazon.com/Apple-Cancellation-Tran...	2025-08-21 11:14:29	Phones	0.00
4	Apple AirTag 4 Pack. Keep Track of and find Yo...	4.8	28988.0	10000.0	72.74	72.74	No Badge	Organic	No Coupon	NaN	NaN	NaN	https://m.media-amazon.com/images/I/61bMNCeAUA...	https://www.amazon.com/Apple-MX542LL-A-AirTag-...	2025-08-21 11:14:29	Phones	0.00

Initial data inspection

Check the data types to understand the dataset structure.

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 42675 entries, 0 to 42674
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   product_title                        42675 non-null  object
1   product_rating                      41651 non-null  float64
2   total_reviews                      41651 non-null  float64
3   purchased_last_month                32164 non-null  float64
4   discounted_price                   40613 non-null  float64
5   original_price                     40613 non-null  float64
6   is_best_seller                     42675 non-null  object
7   is_sponsored                       42675 non-null  object
8   has_coupon                         42675 non-null  object
9   buy_box_availability                28022 non-null  object
10  delivery_date                      30692 non-null  object
11  sustainability_tags                 3408 non-null   object
12  product_image_url                  42675 non-null  object
13  product_page_url                   40606 non-null  object
14  data_collected_at                 42675 non-null  object
15  product_category                   42675 non-null  object
16  discount_percentage                40613 non-null  float64
dtypes: float64(6), object(11)
memory usage: 5.5+ MB
```

```
df['is_best_seller'].unique()
```

```
array(['No Badge', 'Best Seller', "Amazon's", 'Limited time deal',
      'Save 30%', 'Save 12%', 'Save 9%', 'Save 17%', 'Save 77%',
      'Ends in', 'Save 10%', 'Save 18%'], dtype=object)
```

```
df['product_category'].unique()
```

```
array(['Phones', 'Laptops', 'Other Electronics', 'Cameras', 'Storage',
      'Printers & Scanners', 'TV & Display', 'Power & Batteries',
      'Chargers & Cables', 'Headphones', 'Networking', 'Speakers',
      'Wearables', 'Gaming', 'Smart Home'], dtype=object)
```

```
df.columns
df = df.drop(columns = ['product_image_url', 'product_page_url', 'data_collected_at', 'delivery_date',
                        'sustainability_tags'])
```

Check missing values

Define a function to summarize missing values across columns.

```
def missing_values(df):
    for i in df.columns:
        print(f"{i}: {df[i].isnull().sum()}, {df[i].isnull().sum() / len(df) * 100}")
missing_values(df)
```

```
product_title: 0, 0.0
product_rating: 1024, 2.3995313415348565
total_reviews: 1024, 2.3995313415348565
purchased_last_month: 10511, 24.630345635618042
discounted_price: 2062, 4.83186877562976
original_price: 2062, 4.83186877562976
is_best_seller: 0, 0.0
is_sponsored: 0, 0.0
has_coupon: 0, 0.0
buy_box_availability: 14653, 34.33626244874048
product_category: 0, 0.0
discount_percentage: 2062, 4.83186877562976
```

Data cleaning

Deal with rows with missing values in key columns to ensure data quality.

```
df = df.dropna(subset = ['product_rating', 'purchased_last_month', 'discounted_price'])
df['buy_box_availability'] = df['buy_box_availability'].fillna(False)

missing_values(df)
```

```
product_title: 0, 0.0
product_rating: 0, 0.0
total_reviews: 0, 0.0
purchased_last_month: 0, 0.0
discounted_price: 0, 0.0
original_price: 0, 0.0
is_best_seller: 0, 0.0
is_sponsored: 0, 0.0
has_coupon: 0, 0.0
buy_box_availability: 0, 0.0
product_category: 0, 0.0
discount_percentage: 0, 0.0
```

Feature engineering

Create new features such as coupon and badge indicators to capture their impact on purchases.

```
df['badge_group'] = (df['is_best_seller'] != 'No Badge').astype(int)
df['badge_group']
```

```
0      0
1      0
2      0
3      1
4      0
..
42670   0
42671   0
42672   1
42673   0
42674   0
Name: badge_group, Length: 30228, dtype: int64
```

```
df['has_coupon'] = df['has_coupon'].apply(lambda x: 'Coupon' if isinstance(x, str) and 'with coupon' in
x.lower() else 'No Coupon')

df['is_best_seller'] = df['is_best_seller'].apply(lambda x: 'Best Seller' if isinstance(x, str) and 'best
seller' in x.lower() else 'No Badge')

df = pd.get_dummies(df, columns=['is_best_seller', 'is_sponsored', 'has_coupon', 'buy_box_availability'],
drop_first=True)

df.head()
```

	product_title	product_rating	total_reviews	purchased_last_month	discounted_price	original_price	product_category	discount_percentage	badge_group	is_best_seller_No Badge	is_sponsored_Sponsored	has_coupon_No Coupon	buy_box_availability_Add to cart
0	BOYA BOYALINK 2 Wireless Lavalier Microphone f...	4.6	375.0	300.0	89.68	159.00	Phones	43.60	0	True	True	False	True
1	LISEN USB C to Lightning Cable, 240W 4 in 1 Ch...	4.3	2457.0	6000.0	9.99	15.99	Laptops	37.52	0	True	True	True	True
2	DJI Mic 2 (2 TX + 1 RX + Charging Case), Wirel...	4.6	3044.0	2000.0	314.00	349.00	Laptops	10.03	0	True	True	True	True
3	Apple AirPods Pro 2 Wireless Earbuds, Active N...	4.6	35882.0	10000.0	162.24	162.24	Phones	0.00	1	False	False	True	False
4	Apple AirTag 4 Pack. Keep Track of and find Yo...	4.8	28988.0	10000.0	72.74	72.74	Phones	0.00	0	True	False	True	False

```
df.columns
```

```
Index(['product_title', 'product_rating', 'total_reviews',  
      'purchased_last_month', 'discounted_price', 'original_price',  
      'product_category', 'discount_percentage', 'badge_group',  
      'is_best_seller_No Badge', 'is_sponsored_Sponsored',  
      'has_coupon_No Coupon', 'buy_box_availability_Add to cart'],  
      dtype='object')
```

Descriptive statistics

Summarize key variables to understand their distribution and central tendencies.

```
# Descriptive Statistics  
main_vars = ["product_rating", "purchased_last_month"]  
cols = main_vars  
  
stats = pd.DataFrame({  
    "count": df[cols].count(),  
    "mean": df[cols].mean(numeric_only=True),  
    "median": df[cols].median(numeric_only=True),  
    "std": df[cols].std(numeric_only=True, ddof=1),  
    "min": df[cols].min(numeric_only=True),  
    "max": df[cols].max(numeric_only=True),  
    "nan_count": df[cols].isna().sum(),  
}).round(3)
```

```
stats
```

	count	mean	median	std	min	max	nan_count
product_rating	30228	4.439	4.5	0.359	2.0	5.0	0
purchased_last_month	30228	1354.570	200.0	6476.654	50.0	100000.0	0

From the statistics of the two variables we can find interesting insights:

1. We have more ratings than purchases. This is due to the fact that a lot of products were not bought last month but still have a rating from past purchases. We can see it from the number of nan_counts.
2. In stars the median is 4.5 and the mean is 4.399. The mean being lower is because of products with super low ratings. With purchased_last_month, the median is 200 and the mean is 1293.665. The mean is higher because there are products with a high numbers of purchases in the last month.
3. We also have products that have no rating.

Data preparation

Select relevant columns and convert variables to numeric format for analysis.

```
# keep only needed cols, drop missing, ensure numeric
df_plot = df[['product_rating', 'purchased_last_month']].copy()
df_plot = df_plot.dropna()
df_plot['product_rating'] = pd.to_numeric(df_plot['product_rating'], errors='coerce')
df_plot['purchased_last_month'] = pd.to_numeric(df_plot['purchased_last_month'], errors='coerce')
df_plot = df_plot.dropna()
```

Rating distribution

Visualize how product ratings are distributed.

```
# Histogram: Distribution of Star Ratings
plt.figure(figsize=(6,4))


# Define bins up to 5.0 only
bins = np.arange(0.5, 5.1, 0.5)

plt.hist(df["product_rating"], bins=bins, color='skyblue', edgecolor='black')

plt.title('Distribution of Product Star Ratings')
plt.xlabel('Star Rating')
plt.ylabel('Number of Products')
plt.grid(axis='y', linestyle='--', alpha=0.7)

# Set ticks only up to 5.0
plt.xticks(np.arange(0.5, 5.1, 0.5))
plt.xlim(0.5, 5.0)

plt.tight_layout()
plt.show()
```



Purchases distribution


Explore the distribution of monthly product purchases.

```
# Histogram: Distribution of Sales (Purchases Last Month)
fig, axes = plt.subplots(1, 2, figsize=(12, 4))

# Left: Full distribution
axes[0].hist(df_plot['purchased_last_month'], bins=30, color='skyblue', edgecolor='black')
axes[0].set_title('Distribution of Purchases')
axes[0].set_xlabel('Number of Purchases')
axes[0].set_ylabel('Number of Products')
axes[0].grid(axis='y', linestyle='--', alpha=0.7)

# Right: Capped distribution (≤ 3000)
axes[1].hist(df_plot.loc[df_plot['purchased_last_month'] <= 3000, 'purchased_last_month'],
             bins=30, color='skyblue', edgecolor='black')
axes[1].set_title('Zoomed Distribution of Purchases (≤ 3,000 Purchases)')
axes[1].set_xlabel('Number of Purchases')
axes[1].set_ylabel('Number of Products')
axes[1].grid(axis='y', linestyle='--', alpha=0.7)

plt.tight_layout()
plt.show()
```




Relationship between rating and purchases

Analyze how product ratings relate to sales performance.


```
# Scatterplot: Star Rating vs. Sales
plt.figure(figsize=(6,4))
plt.scatter(df_plot['product_rating'], df_plot['purchased_last_month'], alpha=0.5, color='skyblue')
plt.title('Star Rating vs. Purchases Last Month')
plt.xlabel('Star Rating')
plt.ylabel('Purchases Last Month')
plt.grid(True, linestyle='--', alpha=0.6)

# Set x-axis from 0.0 to 5.0, with ticks every 0.5
plt.xlim(0, 5.0)
plt.xticks(np.arange(0, 5.1, 0.5))

plt.tight_layout()
plt.show()
```



Grouping ratings

Group ratings into bins to simplify analysis and visualization.

```
# Group to nearest 0.5-star instead of integer
df_plot['rating_group'] = (df_plot['product_rating'] * 2).round() / 2
df_plot['rating_group'] = df_plot['rating_group'].clip(1, 5)

# Define the 0.5-step groups
groups = np.arange(1.0, 5.1, 0.5) # 1.0, 1.5, ..., 5.0

# Build data + labels (skip empty groups to avoid empty boxes)
data = [df_plot.loc[df_plot['rating_group'] == g, 'purchased_last_month'] for g in groups]
labels = [str(g) for g, s in zip(groups, data) if not s.empty]
data = [s for s in data if not s.empty]
```

Boxplot analysis

Compare purchases across different rating groups to identify patterns.


```
# Side-by-side boxplots
# Note: since we already know in our data, no orders have ratings below 2
# so that's why our boxplot's x-axis automatically starts from 2

fig, axes = plt.subplots(1, 2, figsize=(12, 4))

# Left: Full range
b1 = axes[0].boxplot(data, patch_artist=True)
for patch in b1['boxes']:
    patch.set(facecolor='skyblue', edgecolor='black')
axes[0].set_title('Sales by Star Rating (Full Range)')
axes[0].set_xlabel('Star Rating')
axes[0].set_ylabel('Purchases Last Month')
axes[0].set_xticklabels(labels)
axes[0].grid(axis='y', linestyle='--', alpha=0.7)

# Right: Zoomed ( $\leq 3,000$ )
b2 = axes[1].boxplot(data, patch_artist=True)
for patch in b2['boxes']:
    patch.set(facecolor='skyblue', edgecolor='black')
axes[1].set_ylim(0, 3000)
axes[1].set_title('Sales by Star Rating ( $\leq 3,000$  Purchases)')
axes[1].set_xlabel('Star Rating')
axes[1].set_ylabel('Purchases Last Month')
axes[1].set_xticklabels(labels)
axes[1].grid(axis='y', linestyle='--', alpha=0.7)

plt.tight_layout()
plt.show()
```



Insights for our x and y variables:

- The distribution of rating is left-skewed. Majority of the ratings are between 4.0 inclusive and 5.0 exclusive. The mode is 4.5.
- The distribution of sales is right-skewed. Majority of the products have been bought less than 500 times in the past month.
- In the start rating vs purchases in the last month, we see that there is a weak positive correlation between both variables.
- In the box plots we can see that the number of outliers increases as the star rating increases. Same happens with the variability of the star rating.

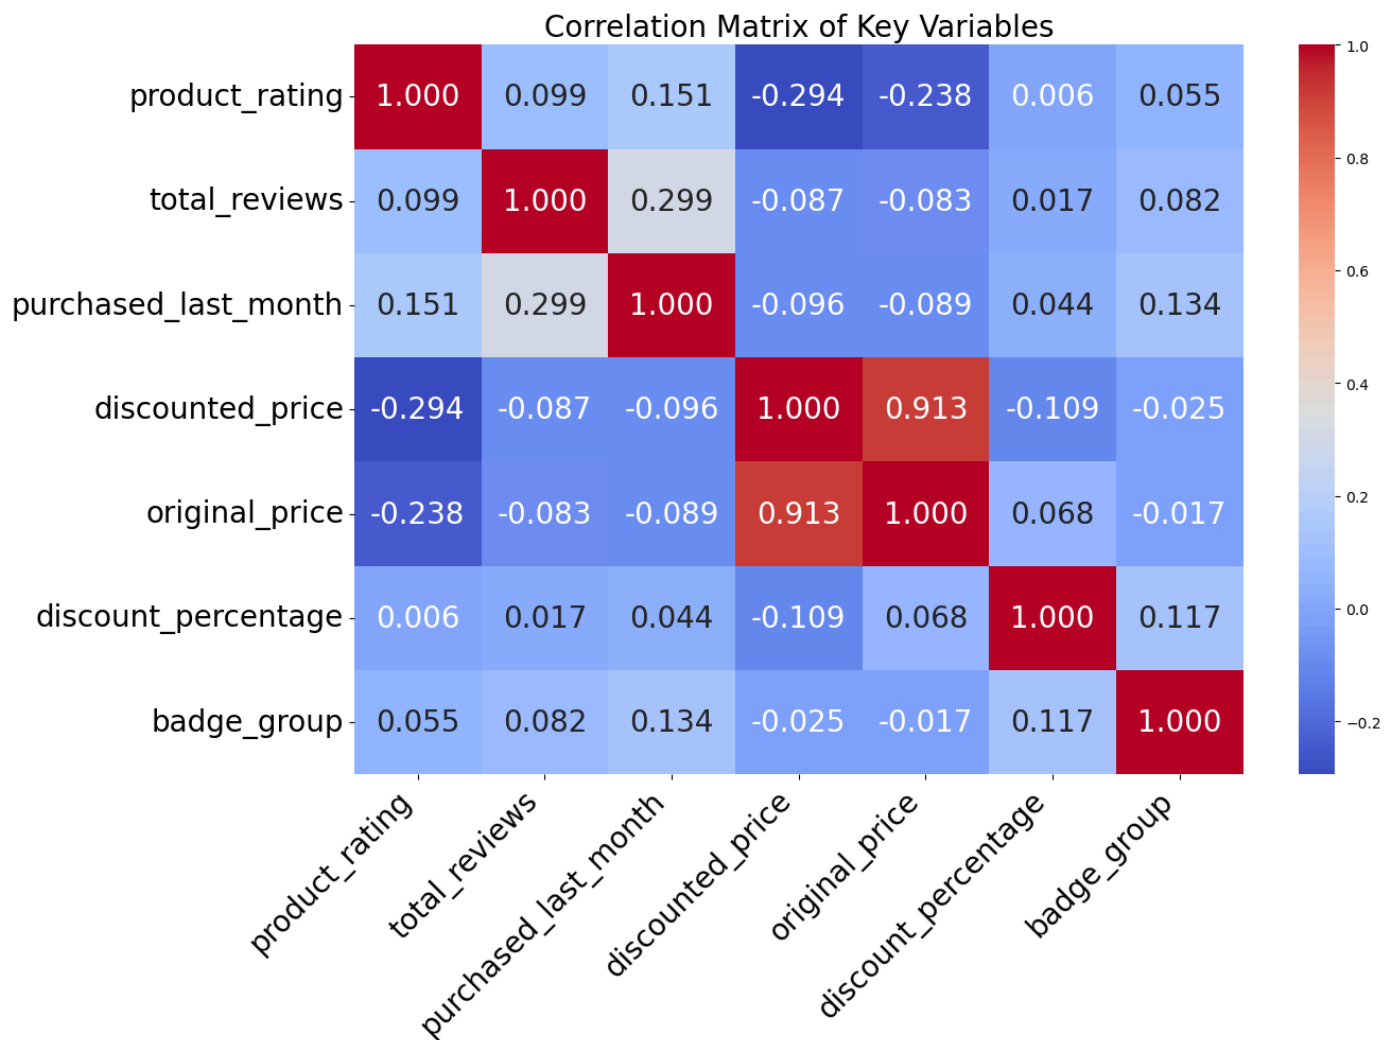
Correlation analysis

Examine relationships between numerical variables using correlation.

```
# Correlation Matrix & correlation coefficient r
corr_vars = df.select_dtypes(include=['int64', 'float64'])
corr_matrix = corr_vars.corr(method='pearson')

plt.figure(figsize=(14,10))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".3f",annot_kws={"size": 20})
plt.title("Correlation Matrix of Key Variables", fontsize=20)
plt.xticks(fontsize=20, rotation=45, ha='right')
plt.yticks(fontsize=20, rotation=0)
plt.tight_layout()
plt.show()

corr_matrix
```



	product_rat ing	total_revie ws	purchased_last_m onth	discounted_p rice	original_pr ice	discount_percen tage	badge_gro up
product_rating	1.000000	0.099372	0.151390	-0.294041	-0.238463	0.006486	0.055282
total_reviews	0.099372	1.000000	0.299240	-0.087026	-0.083030	0.016975	0.082394
purchased_last_m onth	0.151390	0.299240	1.000000	-0.095771	-0.089010	0.044170	0.134451
discounted_price	-0.294041	-0.087026	-0.095771	1.000000	0.913084	-0.108510	-0.025393
original_price	-0.238463	-0.083030	-0.089010	0.913084	1.000000	0.068095	-0.017453

	product_rating	total_reviews	purchased_last_month	discounted_price	original_price	discount_percentage	badge_group
discount_percentage	0.006486	0.016975	0.044170	-0.108510	0.068095	1.000000	0.117411
badge_group	0.055282	0.082394	0.134451	-0.025393	-0.017453	0.117411	1.000000

Confounding variables

Explore additional factors that may influence purchases beyond ratings.

```
# Confounding Variable Plots
df_conf = df.copy()

# Keep only purchases ≤ 3000 for clarity
df_conf = df_conf[df_conf['purchased_last_month'] <= 3000]
```


Product category vs purchases

```
# Product Category vs Purchases
plt.figure(figsize=(8,4))

# Compute average purchases by category
avg_purchases = (
    df_conf.groupby('product_category')['purchased_last_month']
    .mean()
    .sort_values(ascending=False)
)

# Bar plot for average purchases
sns.barplot(
    x=avg_purchases.index,
    y=avg_purchases.values,
    color='skyblue'
)

plt.title('Average Purchases by Product Category ( $\leq 3,000$  Purchases)')
plt.xlabel('Product Category')
plt.ylabel('Average Purchases')
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```




Reviews vs purchases

```
# Total Reviews vs Purchases
plt.figure(figsize=(6,4))

# Scatterplot showing relationship between reviews and purchases
sns.scatterplot(
    x='total_reviews',
    y='purchased_last_month',
    data=df_conf,
    color='skyblue',
    alpha=0.5,
    edgecolor='black',
    linewidth=0.3
)

plt.title('Purchases vs Total Reviews ( $\leq 3,000$  Purchases)')
plt.xlabel('Total Reviews')
plt.ylabel('Purchases ( $\leq 3,000$ )')
plt.grid(True, linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```




Badge effect on purchases

```
# Badge Group vs Purchases
plt.figure(figsize=(6, 4))

sns.boxplot(
    x='badge_group',
    y='purchased_last_month',
    data=df_conf,
    color='skyblue'
)

plt.title('Purchases by Badge Group ( $\leq 3,000$  Purchases)')
plt.xlabel('Badge Group')
plt.ylabel('Purchases ( $\leq 3,000$ )')
plt.xticks([0, 1], ['(0) No Badge', '(1) Has Badge']) # label 0/1 nicely
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



Three confounding variables:

1. Product Category – Different categories attract different levels of consumer demand; certain types (like electronics or toys) naturally sell more, making category a potential confounding factor affecting both rating and sales.
2. Total Reviews - Products with more reviews often appear more trustworthy or popular, which can increase sales while also correlating with higher ratings, creating overlap in influence.
3. Badge (No Badge, Best Seller, Amazon's Choice, Discount Badge) – Official badges like Best Seller signal product quality and popularity, directly boosting purchases and potentially influencing ratings, thus confounding the relationship between ratings and sales.

Prepare data for modeling

Convert categorical variables into numerical format using dummy encoding.

Then, select numerical features and prepare the dataset for modeling by removing the target variable and adding a constant term.

Finally, we calculate the Variance Inflation Factor (VIF) to detect multicollinearity between variables. High VIF values indicate that some features are highly correlated and may affect model stability.

```
# convert category to numeric
df = pd.get_dummies(df, columns=['product_category'], drop_first=True)
```

```
# keep numeric features
X_num = df.select_dtypes(include=[float, int])
X_num = X_num.drop(columns=['product_rating'])
X_num = add_constant(X_num)

vif_data = pd.DataFrame()
vif_data['Feature'] = X_num.columns
vif_data['VIF'] = [variance_inflation_factor(X_num.values, i) for i in range(X_num.shape[1])]

print(vif_data)
```

	Feature	VIF
0	const	1.983264
1	total_reviews	1.104747
2	purchased_last_month	1.119597
3	discounted_price	7.310382
4	original_price	7.258401
5	discount_percentage	1.237907
6	badge_group	1.034763

```
df.columns
```

```
Index(['product_title', 'product_rating', 'total_reviews',
      'purchased_last_month', 'discounted_price', 'original_price',
      'discount_percentage', 'badge_group', 'is_best_seller_No Badge',
      'is_sponsored_Sponsored', 'has_coupon_No Coupon',
      'buy_box_availability_Add to cart',
      'product_category_Chargers & Cables', 'product_category_Gaming',
      'product_category_Headphones', 'product_category_Laptops',
      'product_category_Networking', 'product_category_Other Electronics',
      'product_category_Phones', 'product_category_Power & Batteries',
      'product_category_Printers & Scanners', 'product_category_Smart Home',
      'product_category_Speakers', 'product_category_Storage',
      'product_category_TV & Display', 'product_category_Wearables'],
      dtype='object')
```

Simple linear regression (rating vs purchases)

We start with a simple linear regression model using only product rating to predict monthly purchases.

This helps us understand the direct relationship between rating and sales before adding more variables.

```

iX = df[['product_rating']]

y = df['purchased_last_month']

for col in iX.select_dtypes(include='bool').columns:
    iX[col] = iX[col].astype(int)

X = sm.add_constant(iX)

model = sm.OLS(y, X).fit()

print(model.summary())

rmse = np.sqrt((model.resid) ** 2).mean()
print(f"RMSE: {rmse:.4f}")

plt.figure(figsize=(8,5))
plt.scatter(iX['product_rating'], y, alpha=0.7, label="Observed data")
plt.plot(iX['product_rating'], model.fittedvalues, color='red', linewidth=2, label="Regression line")
plt.xlabel("Product Rating")
plt.ylabel("Purchased Last Month")
plt.title("Linear Regression Fit")
plt.legend()
plt.show()

```

```


                        OLS Regression Results
=====
Dep. Variable:    purchased_last_month    R-squared:                0.023
Model:                OLS    Adj. R-squared:            0.023
Method:            Least Squares    F-statistic:            709.0
Date:                Mon, 23 Mar 2026    Prob (F-statistic):      2.00e-154
Time:                02:46:49    Log-Likelihood:         -3.0782e+05
No. Observations:    30228    AIC:                    6.156e+05
Df Residuals:        30226    BIC:                    6.157e+05
Df Model:            1
Covariance Type:    nonrobust
=====
                        coef    std err          t      P>|t|      [0.025    0.975]
-----
const            -1.077e+04    456.801    -23.575     0.000    -1.17e+04    -9873.738
product_rating    2731.3088     102.577     26.627     0.000     2530.254     2932.363
=====
Omnibus:            47656.888    Durbin-Watson:           1.829
Prob(Omnibus):      0.000    Jarque-Bera (JB):        22588615.189
Skew:               10.371    Prob(JB):                 0.00
Kurtosis:           135.304    Cond. No.                  58.0
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
RMSE: 6401.8985

Linear Regression Fit



```

# Filter purchases <= 3000 for visualization
df_plot = df[df['purchased_last_month'] <= 3000]

# Prepare data for regression
X = sm.add_constant(df['product_rating'])
y = df['purchased_last_month']
model = sm.OLS(y, X).fit()

# Predictions for regression line
pred_x = np.linspace(df['product_rating'].min(), df['product_rating'].max(), 100)
pred_x = sm.add_constant(pred_x)
pred_y = model.predict(pred_x)


# Plot
plt.figure(figsize=(8,5))
sns.scatterplot(
    data=df_plot,
    x='product_rating',
    y='purchased_last_month',
    alpha=0.5
)

# Regression line
plt.plot(pred_x, pred_y, color='red', linestyle='--', label='Regression Line')

# Limit y-axis to <= 3000 to zoom in and see the trend
plt.ylim(0, 3000)

plt.xlabel("Star Rating")
plt.ylabel("Purchases Last Month (≤ 3000)")
plt.title("Star Rating vs Purchases (Regression Line)")
plt.legend()
plt.show()

```



The regression results show that product rating has a positive and statistically significant relationship with purchases. Higher-rated products tend to have more monthly sales.

However, the R-squared value is very low, indicating that rating alone explains only a small portion of the variation in purchases. This suggests that other factors (such as reviews or product category) also play an important role.

Multiple linear regression

We build a multiple regression model including rating, reviews, discounts, badges, sponsorship, coupons, and product categories.

This helps us understand how multiple factors jointly influence product purchases.

```
# Interaction features
df['int_pr_s'] = df['product_rating'] * df['is_sponsored_Sponsored']
```

```

iX = df[['product_rating', 'discount_percentage', 'total_reviews', 'int_pr_s', 'badge_group',
'is_sponsored_Sponsored', 'has_coupon_No Coupon',
'product_category_Chargers & Cables', 'product_category_Gaming',
'product_category_Headphones', 'product_category_Laptops',
'product_category_Networking', 'product_category_Other Electronics',
'product_category_Phones', 'product_category_Power & Batteries',
'product_category_Printers & Scanners', 'product_category_Smart Home',
'product_category_Speakers', 'product_category_Storage',
'product_category_TV & Display', 'product_category_Wearables']]

y = df['purchased_last_month']

for col in iX.select_dtypes(include='bool').columns:
    iX[col] = iX[col].astype(int)

X = sm.add_constant(iX)

model = sm.OLS(y, X).fit()

print(model.summary())

rmse = np.sqrt(((model.resid) ** 2).mean())
print(f"RMSE: {rmse:.4f}")

plt.figure(figsize=(8,5))
plt.scatter(model.fittedvalues, model.resid, alpha=0.7)
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel("Fitted Values")
plt.ylabel("Residuals")
plt.title("Residuals vs Fitted Values")
plt.show()

```

OLS Regression Results

```

=====
Dep. Variable:    purchased_last_month    R-squared:                0.307
Model:                OLS                Adj. R-squared:         0.307
Method:             Least Squares        F-statistic:            638.5
Date:                Mon, 23 Mar 2026     Prob (F-statistic):      0.00
Time:                02:46:49             Log-Likelihood:         -3.0262e+05
No. Observations:    30228               AIC:                   6.053e+05
Df Residuals:        30206               BIC:                   6.055e+05
Df Model:            21
Covariance Type:     nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-1263.5325	488.658	-2.586	0.010	-2221.323	-305.742
product_rating	-79.3979	96.270	-0.825	0.410	-268.091	109.295
discount_percentage	17.5944	2.361	7.453	0.000	12.967	22.222
total_reviews	0.1108	0.002	52.215	0.000	0.107	0.115
int_pr_s	9768.3125	320.296	30.498	0.000	9140.518	1.04e+04
badge_group	1478.8506	141.114	10.480	0.000	1202.261	1755.440
is_sponsored_Sponsored	-3.925e+04	1447.015	-27.125	0.000	-4.21e+04	-3.64e+04
has_coupon_No Coupon	1068.8974	177.449	6.024	0.000	721.090	1416.704
product_category_Chargers & Cables	231.0020	176.825	1.306	0.191	-115.582	577.586
product_category_Gaming	195.9813	277.361	0.707	0.480	-347.658	739.620
product_category_Headphones	318.1925	214.072	1.486	0.137	-101.397	737.782
product_category_Laptops	394.5438	129.677	3.043	0.002	140.371	648.716
product_category_Networking	366.7980	219.873	1.668	0.095	-64.162	797.758
product_category_Other Electronics	667.9808	131.572	5.077	0.000	410.093	925.868
product_category_Phones	-102.6996	136.761	-0.751	0.453	-370.757	165.358
product_category_Power & Batteries	4997.1473	164.870	30.310	0.000	4673.995	5320.300
product_category_Printers & Scanners	482.3145	213.219	2.262	0.024	64.396	900.233

product_category_Smart Home	-2.0858	358.193	-0.006	0.995	-704.159	699.988
product_category_Speakers	372.6366	205.785	1.811	0.070	-30.710	775.983
product_category_Storage	-636.1585	206.361	-3.083	0.002	-1040.634	-231.683
product_category_TV & Display	117.7784	182.356	0.646	0.518	-239.648	475.205
product_category_Wearables	1615.6016	282.058	5.728	0.000	1062.755	2168.448

```
=====
Omnibus:                44380.242    Durbin-Watson:                1.981
Prob(Omnibus):           0.000    Jarque-Bera (JB):            21122474.112
Skew:                    8.914    Prob(JB):                     0.00
Kurtosis:                131.268    Cond. No.                     7.34e+05
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 7.34e+05. This might indicate that there are strong multicollinearity or other numerical problems.
 RMSE: 5389.7914

```
/var/folders/2s/2f5vr80n5v3dkqq95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
```

```
/var/folders/2s/2f5vr80n5v3dkqq95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
```

```
/var/folders/2s/2f5vr80n5v3dkqq95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
```

```
/var/folders/2s/2f5vr80n5v3dkqq95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
```

```
/var/folders/2s/2f5vr80n5v3dkqq95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
```

```
/var/folders/2s/2f5vr80n5v3dkqq95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
```

```
/var/folders/2s/2f5vr80n5v3dkqq95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
/var/folders/2s/2f5vr80n5v3dkqg95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
/var/folders/2s/2f5vr80n5v3dkqg95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
/var/folders/2s/2f5vr80n5v3dkqg95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
/var/folders/2s/2f5vr80n5v3dkqg95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
/var/folders/2s/2f5vr80n5v3dkqg95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
/var/folders/2s/2f5vr80n5v3dkqg95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
/var/folders/2s/2f5vr80n5v3dkqg95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy


```
iX[col] = iX[col].astype(int)
/var/folders/2s/2f5vr80n5v3dkqg95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
iX[col] = iX[col].astype(int)
/var/folders/2s/2f5vr80n5v3dkqg95t3s_fcm0000gn/T/ipykernel_44216/4164450823.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas->


```
ix[col] = ix[col].astype(int)
```



The multiple regression model shows a significant improvement compared to the simple model ($R^2 = 0.307$), indicating that including additional variables helps better explain product purchases.

Several factors have strong and significant effects:

- **Total reviews** and **discount percentage** positively impact purchases, suggesting that social proof and promotions drive sales.
- **Sponsored products** and their interaction with rating (`int_pr_s`) have a large influence, indicating that visibility and marketing play a major role.
- **Product categories** also show varying effects, meaning demand differs across categories.

Interestingly, **product rating is no longer significant** when other variables are included. This suggests that rating alone is not a key driver of purchases once other factors are considered.